

CPSC 304

Introduction to Database Systems

Schema Refinement and Normal Forms

Textbook Reference

Database Management Systems

Reading: Chapter 19. Concentrate on sections 19.1-19.6
(except 19.5.2)



Databases – the continuing saga

- We've learned that databases are wonderful things
- We've learned how to create a *conceptual design* using ER diagrams
- We've learned how to create a *logical design* by turning the ER diagrams into a relational schema including minimizing the data and relations created
- Now we need to *refine* that schema to reduce duplication of information

Learning Goals



- Debate the pros and cons of redundancy in a database.
- Provide examples of update, insertion, and deletion anomalies.
- Given a set of tables and a set of functional dependencies over them, determine all the keys for the tables.
- Show that a table is/isn't in 3NF or BCNF.
- Prove/disprove that a given table decomposition is a lossless join decomposition. Justify why lossless join decompositions are preferred decompositions.
- Decompose a table into a set of tables that are in 3NF, or BCNF.

Imagine that we've created a perfectly good entity for mailing addresses at UBC:




Instance of the schema on the previous slide

Name	Department	Mailing Location
Jessica Wong	Computer Science	201-2366 Main Mall
Rachel Pottinger	Computer Science	201-2366 Main Mall
Joel Friedman	Computer Science	201-2366 Main Mall
Joel Friedman	Math	121-1984 Mathematics Rd
Mark MacLean	Math	121-1984 Mathematics Rd

What are some problems that might happen with this design?
Based solely on intuition, what might be a better design?

Okay, that's bad. But how do I know for sure if departments have only one address?

- Databases allow you to say that one attribute determines another through a *functional dependency*.
- So if Department determines Address but not Name, we say that there's a functional dependency from Department to Address. But Department is NOT a key.
- Another example:
Address(House#, Street, City, Province, PostalCode)




Functional Dependencies (FDs) – technically speaking

- A functional dependency $X \rightarrow Y$ (where X & Y are sets of attributes)

holds if for every legal instance,
for all tuples $t1, t2$:

if $t1.X = t2.X$ *then* $t1.Y = t2.Y$



- Example:
PostalCode \rightarrow City, Province if:
for each possible $t1, t2$, if $t1.PostalCode = t2.PostalCode$ then
 $t1.\{City, Province\} = t2.\{City, Province\}$
- i.e., given two tuples in r , if the X values agree, then the Y values must also agree
- Also can be read as X *determines* Y

Which functional dependencies do we care about?

- A FD is a statement about *all* allowable instances.
 - Must be identified by application semantics
 - Given some instance $r1$ of R , we can check if $r1$ violates some FD f , but we cannot tell if f holds over R !
- We'll concentrate mostly on cases where there's a single attribute on the RHS: (e.g., $\text{PostalCode} \rightarrow \text{Province}$)
- There are boring, *trivial* cases:
 - e.g. $\text{PostalCode}, \text{House\#} \rightarrow \text{PostalCode}$
- We'll concentrate on the non-boring ones

What do we need to know to split apart addresses without losing information?

- FDs tell us when we're storing redundant information
- Reducing redundancy helps eliminate anomalies and save storage space
- We'd like to split apart tables without losing information
- But first, we need to know both what FDs are *explicit* (given) and what FDs are *implicit* (can be derived)
- Among other things, this can help us derive additional keys from the given keys (spare keys are handy in databases, just like in real life – we'll see why shortly)

The Key's the thing



- As a reminder, a key is a *minimal* set of attributes that uniquely identify a relation
 - i.e., a key is a minimal set of attributes that *functionally determines* all the attributes
 - e.g., House#, Street, PostalCode is a key for Address
- A *superkey* for a relation uniquely identifies the relation, but does not have to be minimal
 - i.e.,: $\text{key} \subseteq \text{superkey}$
 - E.g.,:
 - House#, Street, PostalCode is a key and a superkey
 - House#, Street, PostalCode, city is a superkey, but not a key



Functional dependencies & keys all together



- In a **functional dependency**, a set of attributes determines other attributes, e.g., $AB \rightarrow C$, means A and B together determine C
- A **trivial FD** determines what you already have, eg., $AB \rightarrow B$
- A **key** is a minimal set of attributes determining the rest of the attributes of a relation
For example, $R(\underline{\text{House \#}}, \underline{\text{Street}}, \text{City}, \text{Province}, \underline{\text{Postal Code}})$
- A **superkey** is a set of attributes determining the rest of the attributes in the relation, but does NOT have to be minimal (e.g., the key above, or adding in either of City and Province)
- Given a set of (explicit) functional dependencies, we can determine others...



Deriving Additional FDs: the basics

- Given some FDs, we can often infer additional FDs:
 - $studentid \rightarrow city, city \rightarrow acode$ implies $studentid \rightarrow acode$
- An FD fd is implied by a set of FDs F if fd holds whenever all FDs in F hold.
 - closure of F** : the set of all FDs implied by F .
- Armstrong's Axioms (X, Y, Z are sets of attributes):
 - Reflexivity: If $Y \subseteq X$, then $X \rightarrow Y$
e.g., $city, major \rightarrow city$
 - Augmentation: If $X \rightarrow Y$, then $XZ \rightarrow YZ$ for any Z
e.g., if $sid \rightarrow city$, then $sid, major \rightarrow city, major$
 - Transitivity: If $X \rightarrow Y$ and $Y \rightarrow Z$, then $X \rightarrow Z$
 $sid \rightarrow city, city \rightarrow areacode$ implies $sid \rightarrow areacode$
- These are *sound* and *complete* inference rules for FDs.

Deriving Additional FDs: the extended dance remix



- A couple of additional rules (that follow from axioms):
 - **Union:** If $X \rightarrow Y$ and $X \rightarrow Z$, then $X \rightarrow YZ$
e.g., if $\text{sid} \rightarrow \text{acode}$ and $\text{sid} \rightarrow \text{city}$, then $\text{sid} \rightarrow \text{acode,city}$
 - **Decomposition:** If $X \rightarrow YZ$, then $X \rightarrow Y$ and $X \rightarrow Z$
e.g., if $\text{sid} \rightarrow \text{acode,city}$ then $\text{sid} \rightarrow \text{acode}$, and $\text{sid} \rightarrow \text{city}$

Example: Derive union rule from axioms (Reflexivity, Augmentation, & Transitivity)



Union: If $X \rightarrow Y$ and $X \rightarrow Z$, then $X \rightarrow YZ$
e.g., if $sid \rightarrow acode$ and $sid \rightarrow city$, then $sid \rightarrow acode, city$

- | | | |
|----|---------------------|--------------------|
| 1. | $X \rightarrow Y$ | Given |
| 2. | $X \rightarrow Z$ | Given |
| 3. | $X \rightarrow XY$ | 1, augmentation |
| 4. | $XY \rightarrow ZY$ | 2, augmentation |
| 5. | $X \rightarrow ZY$ | 3, 4, transitivity |

Ex: Derive Decomposition from axioms (Reflexivity, Augmentation, & Transitivity)



Decomposition: If $X \rightarrow YZ$, then $X \rightarrow Y$ and $X \rightarrow Z$
e.g., if $sid \rightarrow acode, city$ then $sid \rightarrow acode$, and $sid \rightarrow city$

- | | | |
|----|--------------------|--------------------|
| 1. | $X \rightarrow YZ$ | Given |
| 2. | $YZ \rightarrow Y$ | Reflexivity |
| 3. | $YZ \rightarrow Z$ | Reflexivity |
| 4. | $X \rightarrow Y$ | 1, 2, transitivity |
| 5. | $X \rightarrow Z$ | 1, 3, transitivity |

All 5 Rules

- **Reflexivity**: If $Y \subseteq X$, then $X \rightarrow Y$
e.g., $\text{city, major} \rightarrow \text{city}$
- **Augmentation**: If $X \rightarrow Y$, then $XZ \rightarrow YZ$ for any Z
e.g., if $\text{sid} \rightarrow \text{city}$, then $\text{sid, major} \rightarrow \text{city, major}$
- **Transitivity**: If $X \rightarrow Y$ and $Y \rightarrow Z$, then $X \rightarrow Z$
 $\text{sid} \rightarrow \text{city}$, $\text{city} \rightarrow \text{areacode}$ implies $\text{sid} \rightarrow \text{areacode}$
- **Union**: If $X \rightarrow Y$ and $X \rightarrow Z$, then $X \rightarrow YZ$
e.g., if $\text{sid} \rightarrow \text{acode}$ and $\text{sid} \rightarrow \text{city}$, then $\text{sid} \rightarrow \text{acode, city}$
- **Decomposition**: If $X \rightarrow YZ$, then $X \rightarrow Y$ and $X \rightarrow Z$
e.g., if $\text{sid} \rightarrow \text{acode, city}$ then $\text{sid} \rightarrow \text{acode}$, and $\text{sid} \rightarrow \text{city}$

Example: Supplier-Part DB

- Suppliers supply parts to projects.
 - SupplierPart(sname,city,status,p#,pname,qty)
 - supplier attributes: sname, city, status
 - part attributes: p#, pname
 - supplier-part attributes: qty
- Functional dependencies:
 - fd1: sname \rightarrow city
 - fd2: city \rightarrow status
 - fd3: p# \rightarrow pname
 - fd4: sname, p# \rightarrow qty
- How can we show that (sname, p#) is a key?

Supplier-Part Key: Part 1: Determining all attributes

Show that (sname, p#) is a superkey of
`SupplierPart(sname,city,status,p#,pname,qty)`

Proof has two parts:

a. Show: `sname, p#` is a (super)key

1. `sname, p# → sname, p#`
2. `sname → city`
3. `sname → status`
4. `sname, p# → city, p#`
5. `sname, p# → status, p#`
6. `sname, p# → sname, p#, status`
7. `sname, p# → sname, p#, status, city`
8. `sname, p# → sname, p#, status, city, qty`
9. `sname, p# → sname, p#, status, city, qty, pname`

fd1:	<code>sname → city</code>
fd2:	<code>city → status</code>
fd3:	<code>p# → pname</code>
fd4:	<code>sname, p# → qty</code>

reflex

fd1

2, fd2, trans

2, aug

3, aug

1, 5, union

4, 6, union

7, fd4, union

8, fd3, union

Supplier-Part Key: Part 2: Minimalness

b. Show: $(sname, p\#)$ is a *minimal key* of
 $SupplierPart(sname, city, status, p\#, pname, qty)$

1. $p\#$ does not appear on the RHS of any FD therefore except for $p\#$ itself, nothing else determines $p\#$
3. specifically, $sname \rightarrow p\#$ does not hold
4. therefore, $sname$ is not a key
5. similarly, $p\#$ is not a key

fd1:	$sname \rightarrow city$
fd2:	$city \rightarrow status$
fd3:	$p\# \rightarrow pname$
fd4:	$sname, p\# \rightarrow qty$

Specifically, for combinations of $sname$ and $p\#$:

- $sname, p\# \rightarrow sname, p\#, city, status, pname, qty$
- $sname \rightarrow sname, city, status$
- $p\# \rightarrow p\#, pname$

Do you, by any chance, have anything less painful?



- Scared you're going to mess up? *Closure* is a fool-proof method of checking FDs and finding implicit FDs.
- Closure for a set of attributes X is denoted X^+
- X^+ includes all attributes of the relation IFF X is a (super)key
- Algorithm for finding Closure of X :

Let Closure = X

Until Closure doesn't change do

if $a_1, \dots, a_n \rightarrow C$ is a FD and $\{a_1, \dots, a_n\} \in \text{Closure}$

Then add C to Closure

fd1:	$\text{sname} \rightarrow \text{city}$
fd2:	$\text{city} \rightarrow \text{status}$
fd3:	$\text{p\#} \rightarrow \text{pname}$
fd4:	$\text{sname}, \text{p\#} \rightarrow \text{qty}$

SupplierPart(sname,city,status,p#,pname,qty)

Ex: $\{\text{sname}, \text{p\#}\}^+ =$

$\{\text{sname}\}^+ =$

$\{\text{p\#}\}^+ =$

--

So seeing if a set of attributes is a key means checking to see if it's closure is all the attributes – pretty simple

Popping back up to our original question...

- Is this really a good design for a table?

Name	Department	Mailing Location
Jessica Wong	Computer Science	201-2366 Main Mall
Rachel Pottinger	Computer Science	201-2366 Main Mall
Laks Lakshmanan	Computer Science	201-2366 Main Mall
Joel Friedman	Computer Science	201-2366 Main Mall
Joel Friedman	Math	121-1984 Mathematics Rd
Mark MacLean	Math	121-1984 Mathematics Rd

- Wouldn't it be nice if there was some rule that said if the amount of redundancy that we had was good?

Approaching Normality

- Role of FDs in detecting redundancy:
 - Consider a relation R with 3 attributes, A B C.
 - **No FDs hold:** There is no redundancy here.
 - **Given $A \rightarrow B$:** Several tuples could have the same A value, and if so, they'll all have the same B value!
- **Normalization:** the process of removing redundancy from data

Normal Forms: Why have one rule when you can have four?

- Provide guidance for table refinement/reducing redundancy.
- Four important normal forms:
 - ***First normal form(1NF)***
 - ***Second normal form (2NF)***
 - ***Third normal form (3NF)***
 - ***Boyce-Codd Normal Form (BCNF)***
- If a relation is in a certain *normal form*, certain problems are avoided/minimized.
- Normal forms can help decide whether decomposition (i.e., splitting tables) will help.



1NF

- Each attribute in a tuple has only one value
 - E.g., for “postal code” you can’t have both V6T 1Z4 and V6S 1W6
- Why do we need it? Codd’s original vision of the relational model allowed multi-valued attributes

1NF

<u>Client ID</u>	<u>Postal Code</u>
1	V6T 1Z4, V6S 1W6

Can't have multiple values
in a single attribute


Normalize
to 1NF


<u>Client ID</u>	<u>Postal Code</u>
1	V6T 1Z4
1	V6S 1W6



2NF

- No partial key dependency
- A relation is in 2NF, if it is in 1NF and for every FD $X \rightarrow Y$ where X is a (minimal) key and Y is a non-key attribute, then no proper subset of X determines Y
- e.g., the address relation is not in 2NF:
 - House#, street, postal_code is a (minimal) key





Boyce-Codd Normal Form (BCNF)

A relation R is in BCNF if:

If $X \rightarrow b$ is a non-trivial dependency in R ,
then X is a superkey for R

(Must be true for every such dependency)

Recall: A dependency is trivial if the LHS contains the RHS, e.g., $\text{City, Province} \rightarrow \text{City}$ is a trivial dependency

In English (though a bit vague):

Whenever a set of attributes of R determine another attribute, it should determine **all** the attributes of R .

What do we want?

Guaranteed freedom from redundancy!

- A relation may be BCNF already!
- Helpful fact: all two attribute relations are in BCNF. Consider relation $S(A,B)$. There are four cases:
 - $A \rightarrow B$ and $B \rightarrow A$
 - Only $A \rightarrow B$
 - Only $B \rightarrow A$
 - No FDs
- Lets look at each one



What if a relation is not in BCNF?

Decomposing a Relation




- A decomposition of R replaces R by two or more relations s.t.:
 - Each new relation contains a subset of the attributes of R (and no attributes not appearing in R), and
 - Every attribute of R appears in at least one new relation.
- Intuitively, decomposing R means storing instances of the relations produced by the decomposition, instead of instances of R.
- E.g., Address(House#, Street, City, Province, Postal Code)
 - Address(House#, Street#, PostalCode),
 - PC(City, Province, PostalCode)

How can we decompose correctly?

A sneak preview: The join

- Definition: $R_1 \bowtie R_2$ is the (natural) join of the two relations; i.e., each tuple of R_1 is concatenated with every tuple in R_2 having the same values on the common attributes.



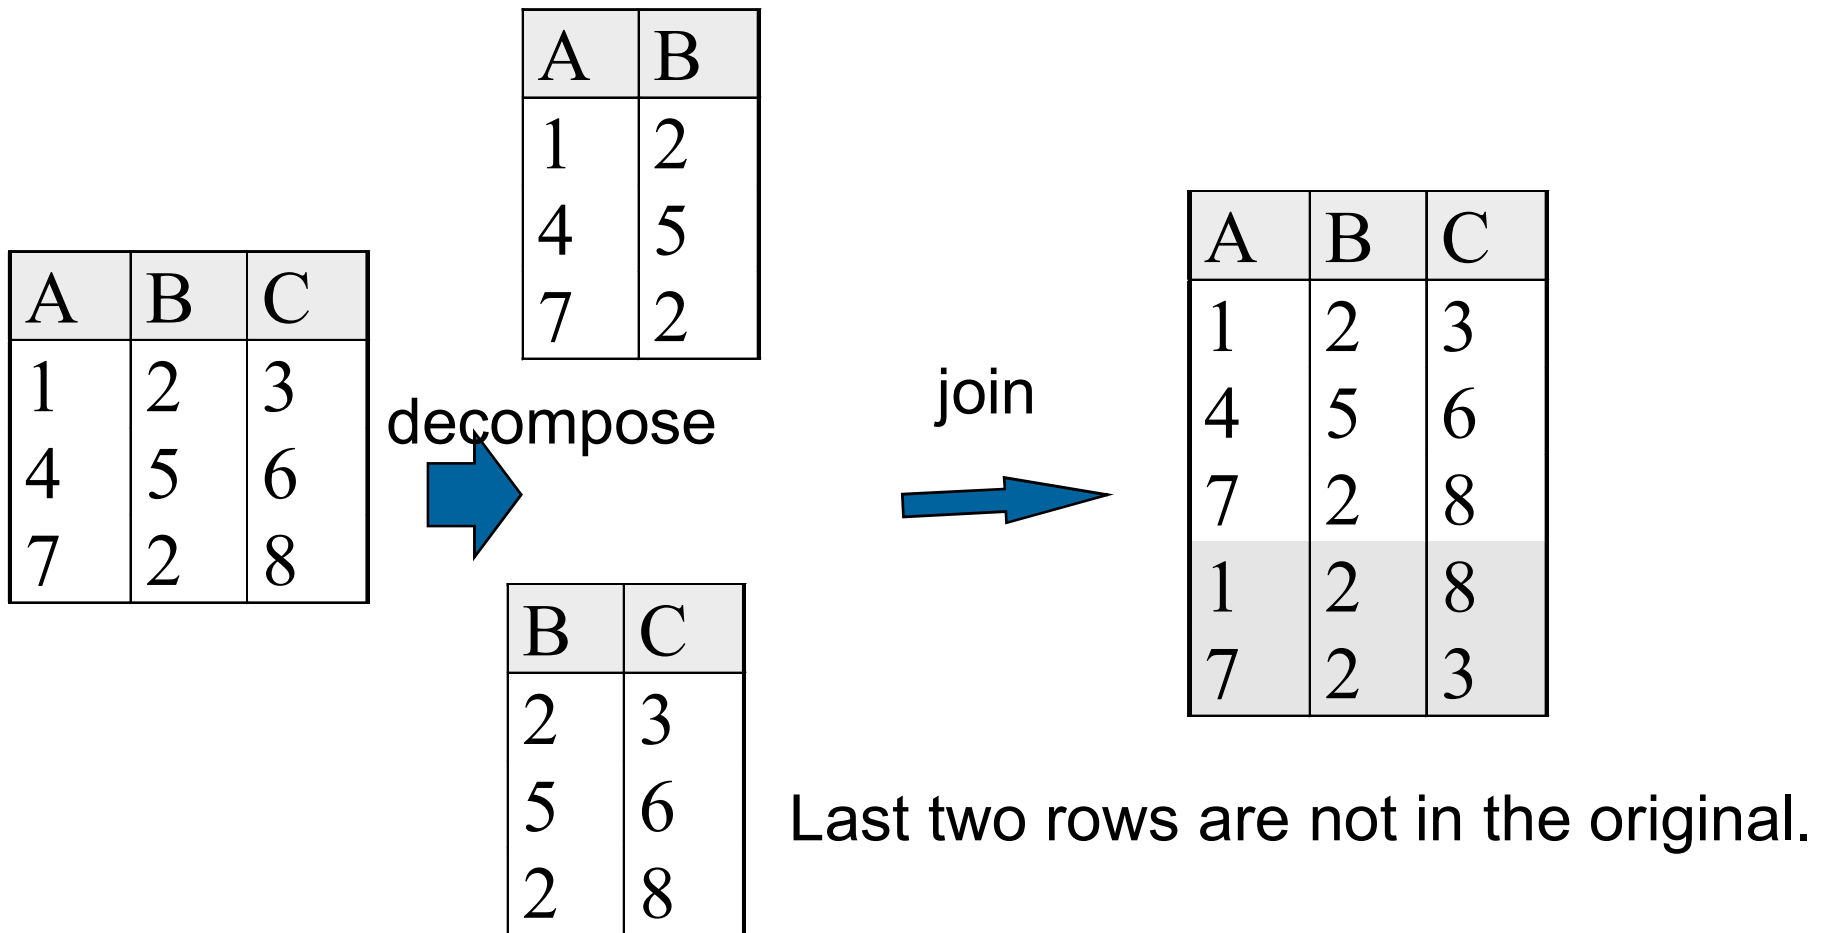
Lossless-Join Decompositions: Definition

Informally: If we break a relation, R , into bits, when we put the bits back together, we should get exactly R back again

Formally: Decomposition of R into X and Y is lossless-join w.r.t. a set of FDs F if, for every instance r that satisfies F :

- If we JOIN the X -part of r with the Y -part of r the result is exactly r
- It is *always* true that r is a subset of the JOIN of its X -part and Y -part, even if the join isn't lossless
- In general, the other direction does not hold – you may get back additional information! If it does hold, the decomposition is a lossless-join.
- Note: The word loss in lossless refers to loss of information, not to loss of tuples. In fact, for “loss of information” a better way to refer to it might be “addition of spurious information”.


Example Lossy-Join Decomposition



All decompositions used to resolve redundancy *must* be lossless!

How do we decompose into BCNF losslessly?

- Let R be a relation with attributes A , and FD be a set of FDs on R s.t. all FDs determine a single attribute
- 1. Pick any $f \in FD$ that violates BCNF of the form $X \rightarrow b$
- 2. Decompose R into two relations: $R_1(A-b)$ & $R_2(X \cup b)$
- Recurse on R_1 and R_2 using FD



Pictorally:

Note: answer may vary depending on order you choose.
That's okay

After you decompose, how do you know which FDs apply?

- Take the closure of the attributes using *all* FDs
- For an FD $X \rightarrow b$ that holds in the original relation, if the decomposed relation S contains $\{X \cup b\}$, then the FD holds for S :
- For example. Consider the relation $S(A, B, C, D, E)$ with functional dependencies $AB \rightarrow C$, $CD \rightarrow E$, $C \rightarrow D$.
- Does $AB \rightarrow E$ hold on S ?
 - First check if A , B and E are all in S ? **They are**
 - Find $AB^+ = \mathbf{ABCDE}$
 - Then yes $AB \rightarrow E$ does hold in S .
- Does $BDE \rightarrow C$ hold? **No**
 - $BDE^+ = BDE$

This BCNF stuff is great and easy!

- Guaranteed that there will be no redundancy of data
- Easy to understand (just look for superkeys)
- Easy to do.
- So what is the main problem with BCNF?
 - For one thing, BCNF may not preserve all dependencies

An illustrative BCNF example

Unit	Company	Product

Unit \rightarrow Company
Company, Product \rightarrow Unit

Is unit a key?

Unit	Company

Unit \rightarrow Company

Unit	Product

We lose the FD: Company, Product \rightarrow Unit !!

So What's the Problem?

<u>Unit</u>	Company
SKYWill	UBC
Team Meat	UBC

Unit	Product
SKYWill	Databases
Team Meat	Databases

Unit → Company

No problem so far. All *local* FD's are satisfied.
Let's put all the data back into a single table again:

Unit	Company	Product
SKYWill	UBC	Databases
Team Meat	UBC	Databases

Violates the FD:

Company, Product → Unit

Review: An illustrative BCNF example

Unit	Company	Product

Unit \rightarrow Company
Company, Product \rightarrow Unit

Is unit a key?

Unit	Company

Unit \rightarrow Company

Unit	Product

We lose the FD: Company, Product \rightarrow Unit !!

Review: So What's the Problem?

<u>Unit</u>	Company
SKYWill	UBC
Team Meat	UBC

Unit	Product
SKYWill	Databases
Team Meat	Databases

Unit → Company

No problem so far. All *local* FD's are satisfied.
Let's put all the data back into a single table again:

Unit	Company	Product
SKYWill	UBC	Databases
Team Meat	UBC	Databases

Violates the FD:

Company, Product → Unit

3NF to the rescue!



A relation R is in 3NF if:

If $X \rightarrow b$ is a non-trivial dependency in R ,
then X is a superkey for R

or b is part of a (minimal) key.

} BCNF

(must be true for every such functional dependency)

Note: b must be part of a **key** *not* part of a **superkey** (if a key exists, *all* attributes are part of a superkey)

Example: $R(\text{Unit}, \text{Company}, \text{Product})$

- $\text{Unit} \rightarrow \text{Company}$
 - $\text{Company}, \text{Product} \rightarrow \text{Unit}$
- Keys: $\{\text{Company}, \text{Product}\}, \{\text{Unit}, \text{Product}\}$

Finding All the Minimal Keys

Find all the minimal keys for R(ABCD) where
 $AB \rightarrow C$ and $C \rightarrow BD$.

Step 1: List all the attributes that **only** appear on the RHS of the FDs.

Left	Middle	Right
		D

Step 2: List all the attributes that **only** ever appear on the LHS of a FD (or do not appear in a FD at all).

Left	Middle	Right
A		D

Finding All the Minimal Keys

Find all the minimal keys for $R(ABCD)$ where $AB \rightarrow C$ and $C \rightarrow BD$.

Step 3: List all the attributes that appear on the LHS and RHS of the FDs.

Left	Middle	Right
A	BC	D

Step 4: Take the closure of the attributes on the LHS.

$$\{A\}^+ = \{A\}$$

Are all of the attributes there? If so, you have found a minimal key. If not, start adding in attributes from the middle column to see if you can determine all the attributes of the relation.

Finding All the Minimal Keys

Find all the minimal keys for R(ABCD) where
 $AB \rightarrow C$ and $C \rightarrow BD$.

Left	Middle	Right
A	BC	D

Step 4.5: Take the closure of the attributes on the LHS.

$$\{A\}^+ = \{A\}$$

$$\{AB\}^+ = \{ABCD\} \leftarrow \text{minimal key}$$

$$\{AC\}^+ = \{ACBD\} \leftarrow \text{minimal key}$$

To decompose to 3NF we rely on the Minimal Cover for a Set of FDs

Goal: Transform FDs to be as small as possible

- Minimal cover G for a set of FDs F:
 - Closure of F = closure of G (i.e., imply the same FDs)
 - Right hand side of each FD in G is a single attribute
 - If we delete an FD in G or delete attributes from an FD in G, the closure changes
- Intuitively, every FD in G is needed, and is “*as small as possible*” in order to get the same closure as F
- e.g., $A \rightarrow B$, $ABCD \rightarrow E$, $EF \rightarrow GH$, $ACDF \rightarrow EG$ has the following minimal cover:
 - $A \rightarrow B$, $ACD \rightarrow E$, $EF \rightarrow G$ and $EF \rightarrow H$

Finding minimal covers of FDs

1. Put FDs in standard form (have only one attribute on RHS)
2. Minimize LHS of each FD
3. Delete Redundant FDs

Example:

$A \rightarrow B$, $ABCD \rightarrow E$, $EF \rightarrow G$, $EF \rightarrow H$, $ACDF \rightarrow EG$

- Replace last rule with
 - $ACDF \rightarrow E$
 - $ACDF \rightarrow G$

Finding minimal covers of FDs

1. Put FDs in standard form (have only one attribute on RHS)
2. Minimize LHS of each FD
3. Delete Redundant FDs

Example:

$A \rightarrow B$, $ABCD \rightarrow E$, $EF \rightarrow G$, $EF \rightarrow H$, $ACDF \rightarrow E$, $ACDF \rightarrow G$

- Can we take anything away from the LHS?
 - $ACD^+ = ABCDE$, (crucially includes B) so remove B from the FD

Finding minimal covers of FDs

1. Put FDs in standard form (have only one attribute on RHS)
2. Minimize LHS of each FD
3. Delete Redundant FDs

Example:

$A \rightarrow B$, $ACD \rightarrow E$, $EF \rightarrow G$, $EF \rightarrow H$, $ACDF \rightarrow E$, $ACDF \rightarrow G$

- Let's find $ACDF^+$ **without** considering the highlighted FDs
 - $ACDF^+ = ACDFEBGH$, so I can remove the highlighted rules
- Final answer: $A \rightarrow B$, $ACD \rightarrow E$, $EF \rightarrow G$, $EF \rightarrow H$

Now we're ready to decompose into 3NF

- We'll cover two methods
- Both methods are
 - Results in relations that do not violate 3NF
 - Lossless (you don't get any additional tuples)
 - Preserve all functional dependencies
- The first one starts by ensuring that the decomposition is lossless and then preserves all functional dependencies
- The second one starts by preserving all functional dependencies and then ensures that the decomposition is lossless

There are other methods that do not ensure both of these properties. Please don't use them!

Decomposition into 3NF method #1: lossless join method

Decomposition into 3NF using the lossless join method:

- Given the FDs F , compute F' : the *minimal cover for F*
- Decompose using F' if violating 3NF similar to how it was done for BCNF – decompose all the way down to BCNF
- After each decomposition identify the set of dependencies N in F' that are not preserved by the decomposition.
- For each $X \rightarrow b$ in N create a relation $R_n(X \cup b)$ and add it to the decomposition

Intuition: first remove redundancy using lossless joins to ensure all results are valid. Then ensure that we maintain all functional dependencies.

3NF example

- Example: $R(ABCDE)$ FD: $AB \rightarrow C, C \rightarrow D$

$ABE^+ = ABCDE$ **only key**

$AB^+ = ABCD$
 $C^+ = CD$

$AB \rightarrow A$


$AB \rightarrow B$

$AB \rightarrow C$

$AB \rightarrow D$


- $AB \rightarrow C$ violates 3NF, so not in 3NF. Decompose

- $R_1(ABC), R_2(ABDE)$



- $AB \rightarrow D$ violates BCNF, so decompose

- $R_3(ABD), R_4(ABE)$



- $R_1(ABC), R_3(ABD), R_4(ABE)$ are now in 3NF
- Are all FDs preserved? We lost $C \rightarrow D$ so add $R_5(CD)$

Decomposition into 3NF Using a Minimal Cover and 3NF Synthesis

- This is an alternative approach to finding a lossless-join, dependency-preserving decomposition into 3NF.
- Consider relation R and its FDs F . To use synthesis to decompose to 3NF:
 1. Find a minimal cover F' .
 2. For each FD $X \rightarrow b$ in F'
Add relation Xb to the decomposition for R .
 3. If there are no relations in the decomposition that contain all of the attributes of a key, add in a relation that contains all the attributes of a key. This preserves lossless joins.

Intuition: first remove redundancy using lossless joins to ensure all results are valid. Then ensure that we maintain all functional dependencies.

Example: Decomposition into 3NF Using a Minimal Cover and 3NF Synthesis

- Suppose we have $R(A,B,C)$ with FDs: $A \rightarrow B$, $C \rightarrow B$.

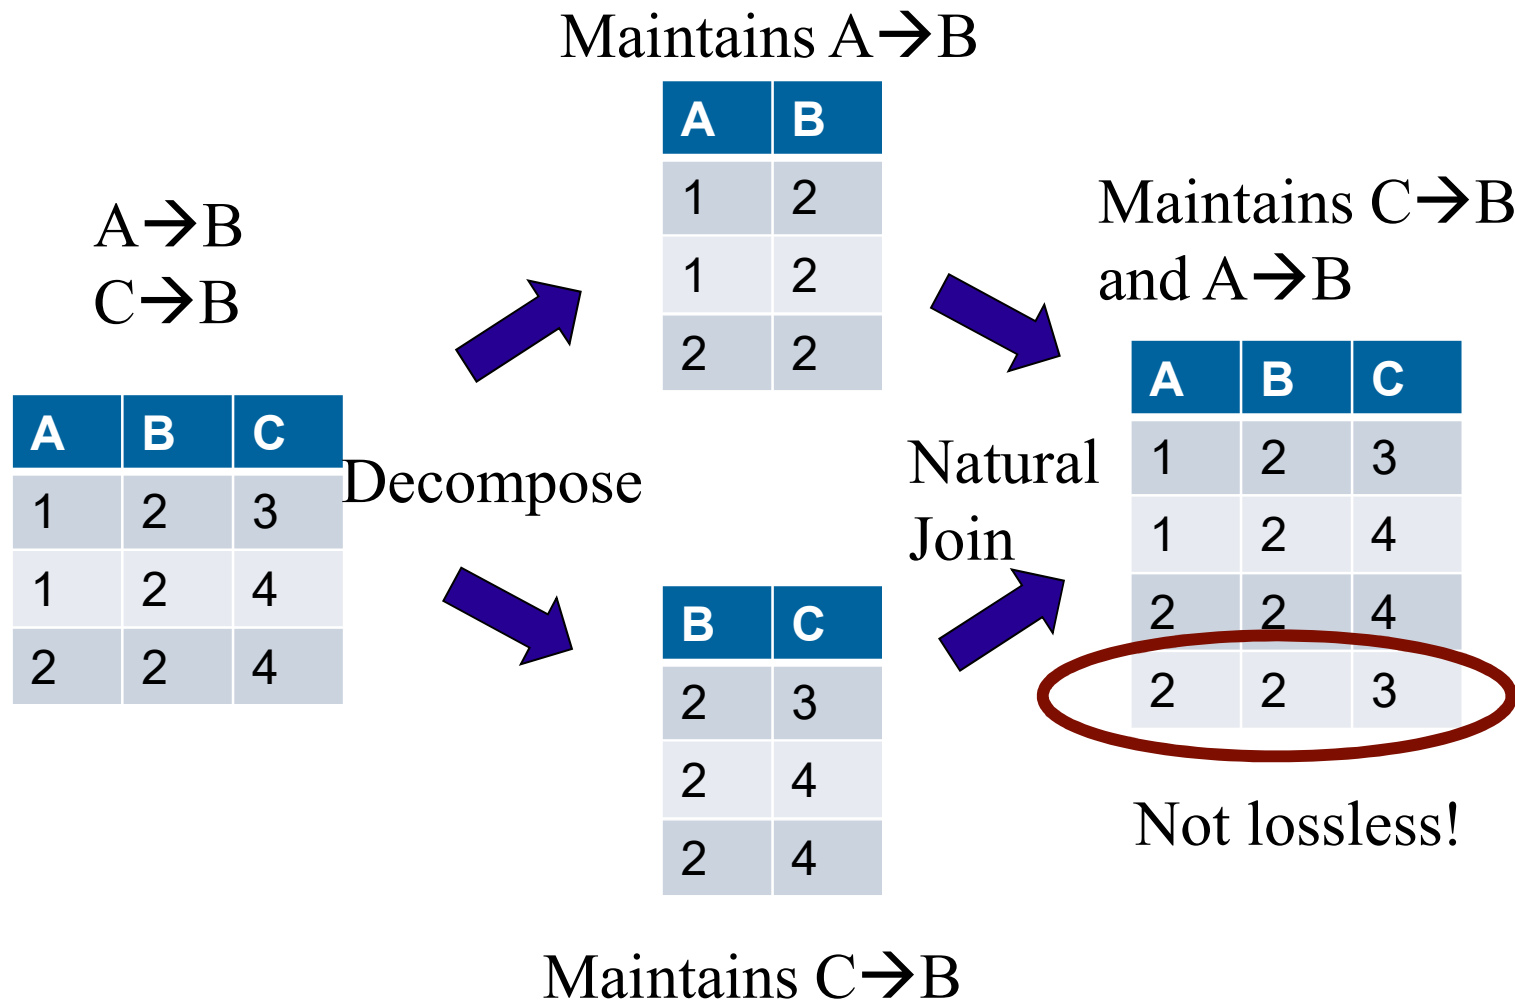
1. Find a minimal cover F' . Already done.
2. For each FD $X \rightarrow b$, add relation Xb to the decomposition for R .

Result: $R_1(A,B)$ and $R_2(B,C)$. Are we done? No.


3. Does it contain a key? What are the keys of R ? AC .
Add $R_3(A,C)$.

- Let's see why we need step #3

Consider the following set of tuples with the previous relation and FDs



Take two!




In decompositions, you can often make some adjustments to make a “better” decomposition

- In this case, R1(ABC), R2(CD), **R3(EFG)**, **R4(EF)**, R5(ABEG) have redundant relations – you don’t need R4 because all information is contained in R3
- You can make the same optimizations in decompositions for BCNF
- Other optimizations exist

Comparing BCNF & 3NF

- BCNF guarantees removal of all anomalies
- 3NF has some anomalies, but preserves all dependencies
- If a relation R is in BCNF it is in 3NF.
- A 3NF relation R may not be in BCNF if all 3 of the following conditions are true:
 - a. R has multiple keys
 - b. Keys are composite (i.e. not single-attributed)
 - c. These keys overlap



Normalization and Design

- Most organizations go to 3NF or better
- If a relation has only 2 attributes, it is automatically in 3NF and BCNF
- Our goal is to use lossless-join for all decompositions and preserve dependencies
- BCNF decomposition is always lossless, but may not preserve dependencies
- Good heuristic:
 - Try to ensure that all relations are in at least 3NF
 - Check for dependency preservation

On the other hand...

Denormalization

- Process of intentionally violating a normal form to gain performance improvements
- Performance improvements:
 - Fewer joins
 - Reduces number of foreign keys
 - Since FD's are often indexed, the number of indexes may be reduced
- Useful if certain queries often require (joined) results, and the queries are frequent enough



Learning Goals Revisited

- Debate the pros and cons of redundancy in a database.
- Provide examples of update, insertion, and deletion anomalies.
- Given a set of tables and a set of functional dependencies over them, determine all the keys for the tables.
- Show that a table is/isn't in 3NF or BCNF.
- Prove/disprove that a given table decomposition is a lossless join decomposition. Justify why lossless join decompositions are preferred decompositions.
- Decompose a table into a set of tables that are in 3NF, or BCNF.